

LinkedIn AI Post Filter

A Chrome/Edge/Brave extension that hides or blurs likely AI-generated posts in your LinkedIn feed.

Install (unpacked, takes ~30 seconds)

1. Download this folder to your computer.
2. Open `chrome://extensions` (or `edge://extensions`, `brave://extensions`).
3. Toggle **Developer mode** on (top right).
4. Click **Load unpacked** and select the `linkedin-ai-filter` folder.
5. Open LinkedIn. Posts get scored as you scroll.
6. Click the extension icon to adjust the threshold or switch between hide/blur.

How it scores a post

Each post gets a score based on linguistic tells. The defaults fire on things like:

- **Buzzword phrases** — "leverage", "delve into", "tapestry", "synergy", "paradigm shift"
- **AI openers** — "In today's fast-paced...", "I'm thrilled to share...", numbered listicles ("5 lessons I learned...")
- **Stock constructions** — "It's not X, it's Y", "This isn't just..."
- **Em-dash density** — three or more em-dashes in one post
- **Emoji bullet structure** — multiple lines starting with an emoji
- **Hashtag stuffing** — eight or more hashtags
- **Staccato formatting** — every sentence on its own short line

Posts at or above your threshold (default: **4**) get flagged.

Tuning

- **Threshold 2–3**: aggressive. Catches more, more false positives.

- **Threshold 4–5:** balanced default. Catches obvious AI slop, leaves human posts alone.
- **Threshold 6+:** only the most flagrant offenders.

Every flagged post has a “Show anyway” button — use that if you spot a false positive.

Honest limitations

- **Heuristic only.** It looks at text patterns, not actual provenance. A real human who happens to write listicles with em-dashes will get caught. A skilled prompter who tells the model to write like a normal person will slip through.
- **English-tuned.** The phrase list is English. Other languages will under-match.
- **DOM-fragile.** LinkedIn changes its HTML. If the selectors in `content.js` break, posts will stop being detected — open `content.js` and update the selectors in `findPosts()` and `getPostText()`.
- **TOS.** Modifying the LinkedIn page client-side for personal use is generally fine, but you’re on your own there.

Customizing the detector

Open `content.js` and edit the `AI_PHRASES` array. Each entry is `[phrase, weight]`. Add your own pet peeves with whatever weight feels right (1 = mild signal, 2–3 = strong signal). Reload the extension on `chrome://extensions` after editing.

If you want stricter “It’s not X, it’s Y” detection, or different listicle patterns, edit `AI_REGEX`.

Possible upgrades (not built in)

- Per-author allowlist (never flag posts from people you trust)
- Send post text to an LLM for second-opinion classification (would need user’s API key + a background script)
- Track flagged-but-overridden posts to tune the threshold automatically